

Tatsiana Sokalava

Time Series Forecasting

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: import requests
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: url = "http://archive.ics.uci.edu/ml/machine-learning-databases/00352/Online%20R
resp = requests.get(url)
with open('test.xlsx', 'wb') as output:
    output.write(resp.content)
```

```
In [4]: df=pd.read_excel('test.xlsx')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

```
In [6]: # The dataframe contains 541,909 observations and 8 features. There are missing
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
```

```
Data columns (total 8 columns):
#      Column      Non-Null Count  Dtype
---  -
0      InvoiceNo    541909 non-null  object
1      StockCode    541909 non-null  object
2      Description   540455 non-null  object
3      Quantity     541909 non-null  int64
4      InvoiceDate    541909 non-null  datetime64[ns]
5      UnitPrice     541909 non-null  float64
6      CustomerID    406829 non-null  float64
7      Country       541909 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

```
In [7]: #Let's shift InvoiceDate as our index and convert County column to a categorical
df['InvoiceDate'] = pd.to_datetime(df['InvoiceDate']).dt.date
```

```
In [8]: df1=df.set_index('InvoiceDate')
```

```
In [9]: df1['Country']=df1['Country'].astype('category')
df1['Country'].value_counts().head(10)
```

```
Out[9]: United Kingdom    495478
Germany                9495
France                 8557
EIRE                   8196
Spain                 2533
Netherlands           2371
Belgium               2069
Switzerland           2002
Portugal              1519
Australia             1259
Name: Country, dtype: int64
```

```
In [10]: df1['InvoiceNo']=df1['InvoiceNo'].astype(str)
```

```
In [11]: # It is evident that there are missing values in CustomerID field.
# Also, this dataset may need to be adjusted for non-positive observations of Un
df1.describe()
```

```
Out[11]:
```

	Quantity	UnitPrice	CustomerID
count	541909.000000	541909.000000	406829.000000
mean	9.552250	4.611114	15287.690570
std	218.081158	96.759853	1713.600303
min	-80995.000000	-11062.060000	12346.000000
25%	1.000000	1.250000	13953.000000
50%	3.000000	2.080000	15152.000000
75%	10.000000	4.130000	16791.000000
max	80995.000000	38970.000000	18287.000000

```
In [12]: df1.head()
```

```
Out[12]:
```

	InvoiceNo	StockCode	Description	Quantity	UnitPrice	CustomerID	Country
--	-----------	-----------	-------------	----------	-----------	------------	---------

InvoiceDate	InvoiceNo	StockCode	Description	Quantity	UnitPrice	CustomerID	Country
InvoiceDate							
2010-12-01	536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	2.55	17850.0	United Kingdom
2010-12-01	536365	71053	WHITE METAL LANTERN	6	3.39	17850.0	United Kingdom
2010-12-01	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2.75	17850.0	United Kingdom
2010-12-01	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	3.39	17850.0	United Kingdom
2010-12-01	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	3.39	17850.0	United Kingdom

In [13]: `df1.tail()`


Out[13]:

InvoiceDate	InvoiceNo	StockCode	Description	Quantity	UnitPrice	CustomerID	Country
2011-12-09	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	0.85	12680.0	France
2011-12-09	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	2.10	12680.0	France
2011-12-09	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	4.15	12680.0	France
2011-12-09	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	4.15	12680.0	France
2011-12-09	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	4.95	12680.0	France

In [14]:

```
# Plot UnitPrice and Quantity
fig, ax = plt.subplots(2, figsize=(12,6))
ax[0].plot(df1['UnitPrice'], label = 'Unit Price')
ax[0].legend()
ax[1].plot(df1['Quantity'], label = 'Quantity')
ax[1].legend()
#In the second graph, there are two symmetrical spikes, which needs to be further
```

Out[14]: `<matplotlib.legend.Legend at 0x7fa43d380190>`



```
In [15]: # Check how many records have zero or less unit price
len(df1[df1.UnitPrice<=0])
```

Out[15]: 2517

```
In [16]: df1[df1.UnitPrice<=0].Description.isna().sum()
```

Out[16]: 1454

```
In [17]: df1[df1.UnitPrice<=0].Description.value_counts()
```

```
Out[17]: check                    159
?                                47
damages                         45
damaged                        43
found                          25
...
HEART OF WICKER LARGE           1
?sold as sets?                  1
CERAMIC HEART FAIRY CAKE MONEY BANK 1
MINI CAKE STAND HANGING STRAWBERY  1
WATERING CAN PINK BUNNY         1
Name: Description, Length: 377, dtype: int64
```

```
In [18]: # Remove it from the dataset
df1=df1[df1.UnitPrice>0]
df1.head()
```

```
Out[18]:
```

	InvoiceNo	StockCode	Description	Quantity	UnitPrice	CustomerID	Country
InvoiceDate							
2010-12-01	536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	2.55	17850.0	United Kingdom
2010-12-01	536365	71053	WHITE METAL LANTERN	6	3.39	17850.0	United Kingdom

2011-12-09	2.08	16446.0	United Kingdom			
2011-12-09	2.08	16446.0	United Kingdom			
	InvoiceNo	StockCode		Description	Quantity	\
InvoiceDate						
2011-01-18	C541433	23166	MEDIUM CERAMIC TOP	STORAGE JAR	-74215	
2011-06-20	C557508	23166	MEDIUM CERAMIC TOP	STORAGE JAR	-240	
2011-08-04	C562375	23166	MEDIUM CERAMIC TOP	STORAGE JAR	-12	
2011-10-12	C570867	23166	MEDIUM CERAMIC TOP	STORAGE JAR	-12	
2011-05-24	C554527	23166	MEDIUM CERAMIC TOP	STORAGE JAR	-9	
...	
2011-05-12	552882	23166	MEDIUM CERAMIC TOP	STORAGE JAR	96	
2011-07-24	561051	23166	MEDIUM CERAMIC TOP	STORAGE JAR	144	
2011-05-18	553607	23166	MEDIUM CERAMIC TOP	STORAGE JAR	240	
2011-07-31	561901	23166	MEDIUM CERAMIC TOP	STORAGE JAR	288	
2011-01-18	541431	23166	MEDIUM CERAMIC TOP	STORAGE JAR	74215	

	UnitPrice	CustomerID	Country
InvoiceDate			
2011-01-18	1.04	12346.0	United Kingdom
2011-06-20	1.04	16684.0	United Kingdom
2011-08-04	1.25	14911.0	EIRE
2011-10-12	1.25	12607.0	USA
2011-05-24	1.04	15251.0	United Kingdom
...
2011-05-12	1.04	14646.0	Netherlands
2011-07-24	1.04	16684.0	United Kingdom
2011-05-18	1.04	16684.0	United Kingdom
2011-07-31	1.25	14156.0	EIRE
2011-01-18	1.04	12346.0	United Kingdom

[260 rows x 7 columns],

	InvoiceNo	StockCode		Description	\
InvoiceDate					
2010-12-02	C536757	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2010-12-06	C537251	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2011-01-05	C540164	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2011-10-24	C572473	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2011-02-28	545217	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
...
2010-12-10	538191	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2010-12-02	536784	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2010-12-15	538998	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2011-10-24	572325	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	
2011-11-22	577822	84347	ROTATING SILVER ANGELS	T-LIGHT HLDR	

	Quantity	UnitPrice	CustomerID	Country
InvoiceDate				
2010-12-02	-9360	0.03	15838.0	United Kingdom
2010-12-06	-9	2.55	NaN	United Kingdom
2011-01-05	-6	2.55	14911.0	EIRE
2011-10-24	-1	2.55	18188.0	United Kingdom
2011-02-28	1	4.96	NaN	United Kingdom
...
2010-12-10	240	1.88	15061.0	United Kingdom
2010-12-02	240	1.88	15061.0	United Kingdom
2010-12-15	480	1.88	15061.0	United Kingdom
2011-10-24	600	1.74	14607.0	United Kingdom
2011-11-22	600	1.74	14607.0	United Kingdom

[475 rows x 7 columns],

	InvoiceNo	StockCode		Description	Quantity
InvoiceDate					
2011-04-18	C550456	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR		-3114
2011-11-22	C577832	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR		-18

2011-12-01	C580131	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	-18
2011-07-28	561658	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	1
2011-06-20	557502	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	1
...
2011-11-23	578125	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	90
2011-11-14	576180	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	180
2011-11-09	575296	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	540
2011-01-11	540815	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114
2011-04-18	550461	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR	3114

	UnitPrice	CustomerID	Country
InvoiceDate			
2011-04-18	2.10	15749.0	United Kingdom
2011-11-22	0.79	18274.0	United Kingdom
2011-12-01	2.10	17576.0	United Kingdom
2011-07-28	2.55	12743.0	Unspecified
2011-06-20	4.96	NaN	United Kingdom
...
2011-11-23	0.79	17511.0	United Kingdom
2011-11-14	0.79	13694.0	United Kingdom
2011-11-09	0.79	16041.0	United Kingdom
2011-01-11	2.10	15749.0	United Kingdom
2011-04-18	2.10	15749.0	United Kingdom

[270 rows x 7 columns],

	InvoiceNo	StockCode	Description	Quantity \
InvoiceDate				
2011-04-18	C550456	21175	GIN + TONIC DIET METAL SIGN	-2000
2011-10-10	C570290	21175	GIN + TONIC DIET METAL SIGN	-12
2011-06-13	C556647	21175	GIN + TONIC DIET METAL SIGN	-12
2011-05-27	C554870	21175	GIN + TONIC DIET METAL SIGN	-3
2010-12-10	C538350	21175	GIN + TONIC DIET METAL SIGN	-1
...
2011-07-14	560080	21175	GIN + TONIC DIET METAL SIGN	192
2011-09-20	567458	21175	GIN + TONIC DIET METAL SIGN	192
2011-11-21	577747	21175	GIN + TONIC DIET METAL SIGN	240
2011-01-11	540815	21175	GIN + TONIC DIET METAL SIGN	2000
2011-04-18	550461	21175	GIN + TONIC DIET METAL SIGN	2000

	UnitPrice	CustomerID	Country
InvoiceDate			
2011-04-18	1.85	15749.0	United Kingdom
2011-10-10	2.55	14665.0	United Kingdom
2011-06-13	2.55	13012.0	United Kingdom
2011-05-27	2.55	15078.0	United Kingdom
2010-12-10	1.85	13798.0	United Kingdom
...
2011-07-14	2.08	17450.0	United Kingdom
2011-09-20	2.66	17450.0	United Kingdom
2011-11-21	2.67	17450.0	United Kingdom
2011-01-11	1.85	15749.0	United Kingdom
2011-04-18	1.69	15749.0	United Kingdom

[825 rows x 7 columns]]


```
In [22]: #By looking at top 5, it is reasonable to assume that the majority of cancelled
# The exception is item 84347, which doesn't have a reasonable explanation witho
# For now, we will not be removing cancelled transaction, since many are voided
# However, we will define a function and remove transactions that have a negativ
```

```
In [23]: def remove_writeoffs(df):
    qnegat=[x for x in df.InvoiceNo if x.startswith('C')]
    return df[(df.Quantity>0) | (df.InvoiceNo.isin(qnegat))].sort_values(by='Qua
df1=remove_writeoffs(df1)
```

```
In [24]: # View daily sales volume by day
df2=df1.groupby(['InvoiceDate'])['Quantity'].sum()
df2.plot()

# The plot doesn't look stationary. While we take a note of it, we will continue
```

Out[24]: <AxesSubplot:xlabel='InvoiceDate'>



```
In [25]: # Check how many quantity of products have been sold online from each country
a = df1['Quantity'].groupby(df1['Country']).agg('sum').sort_values(ascending = False)
print(a)
```

```
Country
United Kingdom    4399357
Netherlands        199552
EIRE               142363
Germany            117446
France             110479
Australia           83345
Sweden              35637
Switzerland         30324
Spain               26813
Japan               25218
Name: Quantity, dtype: int64
```

```
In [26]: # Since the dates range from 12/01/2010-12/09/2011, it looks like there are missing dates
len(df2.index)
```

Out[26]: 305

```
In [27]: # Define a function to reinstate dates
import datetime
def zero_sales(df):
    idx = pd.date_range(df.index.min(),datetime.date(2011,12,9))
    return df.reindex(idx, fill_value=0)
```

```
In [28]: df2=zero_sales(df2)
```

```
In [29]: # Review which days of the week we have no sales reported
pd.DataFrame(df2.index[df2.values==0])[0].dt.day_name().value_counts()
```

Out[29]: Saturday 53


```
Monday      6
Friday      4
Sunday      3
Thursday    1
Wednesday   1
Tuesday     1
Name: 0, dtype: int64
```


```
In [30]: #!pip install holidays
```

```
In [31]: from datetime import date
import holidays
```

```
In [32]: uk_holidays=pd.Series(holidays.UnitedKingdom(years= [2010,2011] ).keys())
```

```
In [33]: # Plot data with holidays
from matplotlib import pylab
def plot_df(data, x, y, title="", xlabel='Date', ylabel='Value', dpi=100):
    plt.figure(figsize=(16,5), dpi=dpi)
    plt.plot(x, y, color='tab:blue')
    plt.gca().set(title=title, xlabel=xlabel, ylabel=ylabel)
    holidays = uk_holidays[uk_holidays.between(data.index.min(), datetime.date(2
    [pylab.axvline(_x, linewidth=0.5, color='r', ls='--') for _x in holidays]
    plt.show()

plot_df(df2, x=df2.index, y=df2.values, title='Daily Sales Volume')
```



```
In [34]: # Define a week between 2011-11-27 and 2011-12-3 and identify the top 3 per object
week=df1[(df1.index>=datetime.date(2011,11,27) ) & (df1.index<=datetime.date(201
top3=week[['StockCode','Quantity']].groupby('StockCode').sum().sort_values(by='Q
top3
```

```
Out[34]:
```

Quantity	
StockCode	
23084	4588
22197	3195
23582	1851

```
In [35]: # Since the dataset has various descriptions per SKU, let's review the Descriptions
df1[df1['StockCode']==top3.index[0]]['Description'].unique()
```

```
Out[35]: array(['RABBIT NIGHT LIGHT'], dtype=object)
```


```
In [36]: df1[df1['StockCode']==top3.index[1]]['Description'].unique()
```

```
Out[36]: array(['POPCORN HOLDER', 'SMALL POPCORN HOLDER'], dtype=object)
```


```
In [37]: df1[df1['StockCode']==top3.index[2]]['Description'].unique()
```

```
Out[37]: array(['VINTAGE DOILY JUMBO BAG RED '], dtype=object)
```

```
In [38]: # Plot the sales quantity timeseries of each item
rabl=df1[df1.StockCode==top3.index[0]].sort_values(by='InvoiceDate')
rabl2=zero_sales(rabl.groupby(['InvoiceDate'])['Quantity'].sum())
plot_df(rabl2, x=rabl2.index, y=rabl2.values, title="RABBIT NIGHT LIGHT", xlabel
```




```
In [39]: # By looking at the RABBIT NIGHT LIGHT, we can infer that there is an increasing
#There is a somewhat seasonality in the observed data. Also, the data may not be
import statsmodels.api as sm
decomposition = sm.tsa.seasonal_decompose(rabl2, model='additive')
from pylab import rcParams
rcParams['figure.figsize'] = 18, 8
fig = decomposition.plot()
plt.show()
```




```
In [40]: # Review autocorrelation function plots
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```


```
fig, ax = plt.subplots(2, figsize=(12,6))
ax[0]=plot_acf(rab12, ax=ax[0], lags=20)
ax[1]=plot_pacf(rab12, ax=ax[1], lags=20)
```




```
In [41]: # Plot the Popcorn Holder data
poph=df1[df1.StockCode==top3.index[1]].sort_values(by='InvoiceDate')
poph2=zero_sales(poph.groupby(['InvoiceDate'])['Quantity'].sum())
plot_df(poph2, x=poph2.index, y=poph2.values, title="POPCORN HOLDER", xlabel='Da
```



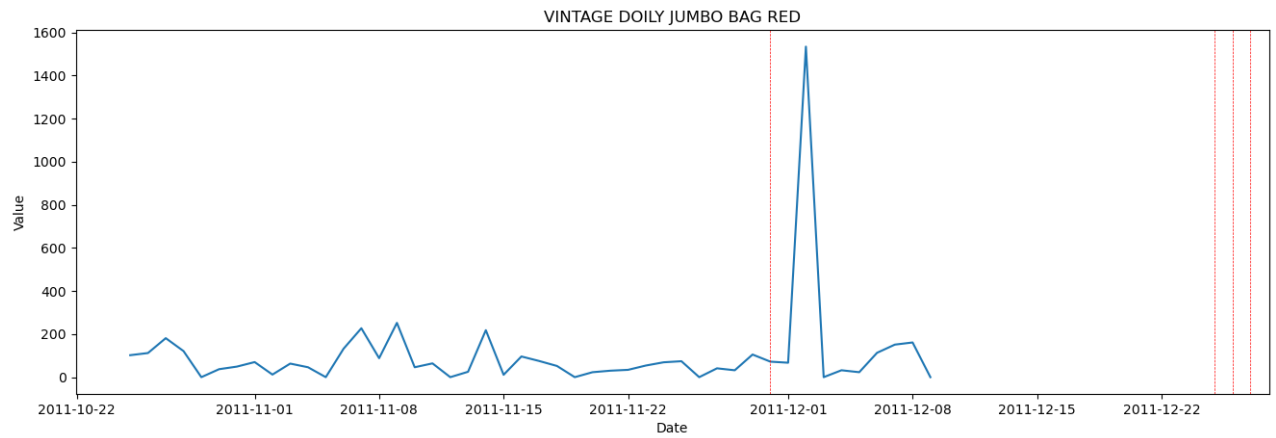
```
In [42]: # By looking at it, we can also infer that there is an unusual spike around June
import statsmodels.api as sm
from pylab import rcParams
decomposition = sm.tsa.seasonal_decompose(poph2, model='additive')
rcParams['figure.figsize'] = 18, 8
fig = decomposition.plot()
plt.show()
```




```
In [43]: fig, ax = plt.subplots(2, figsize=(12,6))
ax[0]=plot_acf(poph2, ax=ax[0], lags=20)
ax[1]=plot_pacf(poph2, ax=ax[1], lags=20)
```




```
In [44]: # Plot sales of the VINTAGE DOILY JUMBO BAG RED
vinb=df1[df1.StockCode==top3.index[2]].sort_values(by='InvoiceDate')
vinb2=zero_sales(vinb.groupby(['InvoiceDate'])['Quantity'].sum())
plot_df(vinb2, x=vinb2.index, y=vinb2.values, title="VINTAGE DOILY JUMBO BAG RED")
```



```
In [45]: # By looking at the VINTAGE DOILY JUMBO BAG RED, we can also infer that there is
decomposition = sm.tsa.seasonal_decompose(vinb2, model='additive')
rcParams['figure.figsize'] = 18, 8
fig = decomposition.plot()
plt.show()
```



```
In [46]: fig, ax = plt.subplots(2, figsize=(12,6))
ax[0]=plot_acf(vinb2, ax=ax[0], lags=20)
ax[1]=plot_pacf(vinb2, ax=ax[1], lags=20)
#The data is stationary and appears to have daily and weekly seasonality along wi
```



In [47]: *#It is evident that three items have a different history of sales: Over a year for RABBIT NIGHT LIGHT, and a month and a half for VINTAGE DOILY JUMBO BAG RED.*

Models

In [48]: `from sklearn.metrics import mean_squared_error`
`from matplotlib import pyplot`

Begin with modeling the POPCORN HOLDER data

`poph2_train=poph2[poph2.index<'2011-11-27']`
`poph2_test=poph2[poph2.index>='2011-11-27']`

#rabl2_train=rabl2[rabl2.index<'2011-11-27']
#rabl2_test=rabl2[rabl2.index>='2011-11-27']

#vinb2_train=vinb2[vinb2.index<'2011-11-27']
#vinb2_test=vinb2[vinb2.index>='2011-11-27']

In [49]: `train_dates, test_dates = poph2_train.index, poph2_test.index`
`train_data, test_data = poph2_train.values, poph2_test.values`

#train_dates, test_dates = rabl2_train.index, rabl2_test.index
#train_data, test_data = rabl2_train.values, rabl2_test.values

#train_dates, test_dates = vinb2_train.index, vinb2_test.index
#train_data, test_data = vinb2_train.values, vinb2_test.values

In [50]: *# Define a dataframe to view the performance of fitted models*
`perform=pd.DataFrame()`
`perform.index.name='Models Popcorn Holder'`
`perform['RMSE']=None`
`perform['Parameters']=None`

In [51]: *# Define a function for RMSE*
`from sklearn.metrics import mean_squared_error`

```
def rmse(actual, predicted):
    rmse=np.sqrt(mean_squared_error(actual,predicted))
    return rmse
```

```
In [52]: # Define a plot function for actual vs predicted:
import plotly.graph_objects as go
def plot_actual_predicted(actual, predicted, model_name):
    fig = go.Figure()
    fig.add_trace(go.Scatter(x=test_dates, y=actual, name = "Expected", line = d
    fig.add_trace(go.Scatter(x=test_dates, y=predicted, name = model_name, line
    fig.show()
```

Model 1: Moving Average

```
In [53]: # Begin selecting the best window size


X=poph2.values
for a in range(1,13):
    window = a
    history = [X[i] for i in range(window)]
    test = [X[i] for i in range(window, (len(X)-14))]
    predictions = list()
    # walk forward over time steps in test
    for t in range(len(test)):
        length = len(history)
        preds = np.mean([history[i] for i in range(length-window,length)])
        obs = test[t]
        predictions.append(preds)
        history.append(obs)
    error = rmse(test, predictions)
    print ('RMSE: %.3f' % round(error), 'Window size: %.3f' % a)
```

```
RMSE: 476.000 Window size: 1.000
RMSE: 409.000 Window size: 2.000
RMSE: 389.000 Window size: 3.000
RMSE: 376.000 Window size: 4.000
RMSE: 370.000 Window size: 5.000
RMSE: 364.000 Window size: 6.000
RMSE: 352.000 Window size: 7.000
RMSE: 350.000 Window size: 8.000
RMSE: 349.000 Window size: 9.000
RMSE: 346.000 Window size: 10.000
RMSE: 346.000 Window size: 11.000
RMSE: 346.000 Window size: 12.000
```

```
In [54]: # Apply a window size of 10 and predict
window = 10
predict=pd.DataFrame(X)[0].rolling(window).mean()[-13:]
observ=pd.DataFrame(X)[0][-13:]
error_ma = rmse(observ, predict)
print('RMSE error: %.3f' % error_ma, 'Window size: %.3f' % window)
```

```
RMSE error: 745.817 Window size: 10.000
```

```
In [55]: plot_actual_predicted(observ, predict, "Moving Average Predictions")
```



```
In [56]: perform.loc['Moving Average', ('RMSE', 'Parameters')] = round(error_ma, 0), 'Window
perform
```

```
Out[56]:
```

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10

Model 2: Exponential Smoothing

```
In [57]: from statsmodels.tsa.holtwinters import ExponentialSmoothing
from sklearn.metrics import mean_squared_error
```

```
In [58]: #!pip install pmdarima
```

```
In [59]: #from statsmodels.tsa.statespace.sarimax import SARIMAX
#from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
In [60]: mod_ex=ExponentialSmoothing(np.array(train_data), seasonal_periods=52, trend='ad
```

```
In [61]: fit=mod_ex.fit()
pred_es=fit.forecast(13)
```

```
In [62]: error_es = rmse(test_data, pred_es)
print('RMSE: %.3f' % error_es)
```



RMSE: 781.167

```
In [63]: perform.loc['Exponential Smoothing', ('RMSE', 'Parameters')] = round(error_es, 'S')
perform
```

```
Out[63]:
```

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10
Exponential Smoothing	781	Seasonal_periods =52

```
In [64]: plot_actual_predicted(test_data, pred_es, "Exponential Smoothing Prediction")
```



Model 3: Arima

```
In [65]: #!pip install pmdarima
from pmdarima.arima import auto_arima
from statsmodels.tsa.stattools import adfuller
```

```
In [66]: # Let's use auto_arima. We know our data shows seasonality and needs differencing
stepwise_model = auto_arima(train_data, start_p=1, start_q=1, max_p=3, max_q=3,
                             trace=True, error_action='ignore', suppress_warnings=
```

Performing stepwise search to minimize aic

```

ARIMA(1,1,1)(0,1,1)[12] : AIC=inf, Time=0.54 sec
ARIMA(0,1,0)(0,1,0)[12] : AIC=5531.538, Time=0.02 sec
ARIMA(1,1,0)(1,1,0)[12] : AIC=5334.731, Time=0.15 sec
ARIMA(0,1,1)(0,1,1)[12] : AIC=inf, Time=0.28 sec
ARIMA(1,1,0)(0,1,0)[12] : AIC=5445.079, Time=0.03 sec
ARIMA(1,1,0)(2,1,0)[12] : AIC=5291.955, Time=0.32 sec
ARIMA(1,1,0)(2,1,1)[12] : AIC=inf, Time=0.94 sec
ARIMA(1,1,0)(1,1,1)[12] : AIC=inf, Time=0.36 sec
ARIMA(0,1,0)(2,1,0)[12] : AIC=5387.740, Time=0.28 sec
ARIMA(2,1,0)(2,1,0)[12] : AIC=5256.161, Time=0.44 sec
ARIMA(2,1,0)(1,1,0)[12] : AIC=5300.342, Time=0.21 sec
ARIMA(2,1,0)(2,1,1)[12] : AIC=inf, Time=1.17 sec
ARIMA(2,1,0)(1,1,1)[12] : AIC=inf, Time=0.53 sec
ARIMA(3,1,0)(2,1,0)[12] : AIC=5235.364, Time=0.48 sec
ARIMA(3,1,0)(1,1,0)[12] : AIC=5282.255, Time=0.27 sec
ARIMA(3,1,0)(2,1,1)[12] : AIC=inf, Time=1.38 sec
ARIMA(3,1,0)(1,1,1)[12] : AIC=inf, Time=0.63 sec
ARIMA(3,1,1)(2,1,0)[12] : AIC=inf, Time=1.28 sec
ARIMA(2,1,1)(2,1,0)[12] : AIC=inf, Time=1.25 sec
ARIMA(3,1,0)(2,1,0)[12] intercept : AIC=5237.368, Time=0.97 sec

```

Best model: ARIMA(3,1,0)(2,1,0)[12]
Total fit time: 11.532 seconds

```

In [67]: stepwise_model.fit(train_data)
         preds = stepwise_model.predict(n_periods=13)
         preds

```

```

Out[67]: array([ 88.78128757, 118.38078031, 25.54930964, 140.30679137,
                585.01505017, 41.36413624, 1147.21044358, 323.36986013,
                156.51130324, 79.41973066, 139.3972984 , 543.33842197,
                3.84710817])

```

```

In [68]: error_ar=rmse(test_data, preds)

         print('RMSE: %.3f' % error_ar)

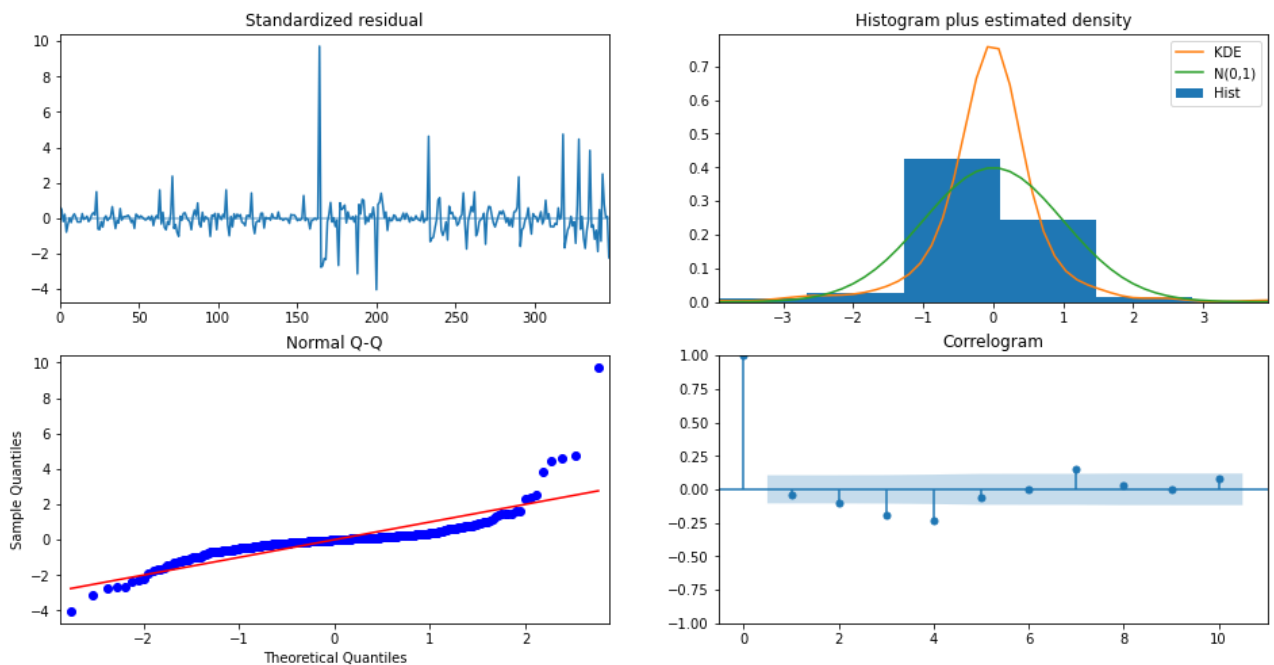
```

RMSE: 909.169


```

In [69]: stepwise_model.plot_diagnostics(figsize=(16, 8))
         plt.show()

```



```
In [70]: plot_actual_predicted(test_data, preds, 'Arima Predictions')
```



```
In [71]: perform.loc['Arima', ('RMSE', 'Parameters')] = round(error_ar), stepwise_model
perform
```

Out[71]:

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10
Exponential Smoothing	781	Seasonal_periods =52
Arima	909	ARIMA(3,1,0)(2,1,0)[12]

Model 4: XGBoost

```
In [72]: #!pip install xgboost
```

```
In [73]: import xgboost as xgb
from xgboost import DMatrix
from pandas import concat
from numpy import asarray
from xgboost import XGBRegressor
```

```

In [74]: # Define functions to apply XGBoost model:
data=poph2.values

# Transform a time series dataset into a supervised learning dataset
def series_to_supervised(data, n_in, n_out=1, dropnan=True):
    df = pd.DataFrame(data)
    cols = list()
    # input sequence (t-n, ... t-1)
    for i in range(n_in, 0, -1):
        cols.append(df.shift(i))
    # forecast out
    for i in range(0, n_out):
        cols.append(df.shift(-i))
    # aggregate
    agg = concat(cols, axis=1)
    # drop rows with NaN values
    if dropnan:
        agg.dropna(inplace=True)
    return agg.values

# Split a dataset into train/test sets
def train_test_split(data, n_test):
    return data[:-n_test, :], data[-n_test:, :]

# Fit an xgboost model and make a one step prediction
def xgboost_forecast(train, testX):
    train = asarray(train)
    # split into input and output columns
    trainX, trainY = train[:, :-1], train[:, -1]
    # fit model
    model = XGBRegressor(objective='reg:squarederror', n_estimators=1000, max_de
min_child_weight=1)
    model.fit(trainX, trainY)
    preds = model.predict(asarray([testX]))
    return preds[0]

# Validation
def walk_forward_test(data, n_test):
    predictions = list()
    train, test = train_test_split(data, n_test)
    # add history with training dataset
    history = [x for x in train]
    # step over each time-step in the test set
    for i in range(len(test)):
        # split test row into input and output columns
        testX, testY = test[i, :-1], test[i, -1]
        # make a prediction
        preds = xgboost_forecast(history, testX)
        predictions.append(preds)
        # add actual observation for the next loop to history
        history.append(test[i])
    # estimate error
    error_xgboost = np.sqrt(mean_squared_error(test[:, -1], predictions))
    return error_xgboost, test[:, -1], predictions

```


```

In [75]: # Transform the time series data into supervised learning
data = series_to_supervised(data, n_in=17)
# Evaluate
error_xgboost, y, preds = walk_forward_test(data, n_test=13)
print('RMSE: %.3f' % error_xgboost)

```

```
# Plot expected vs predicted
plot_actual_predicted(y, preds, "XGBoost Predictions")
```

RMSE: 714.233



```
In [76]: perform.loc['XGBoost', ('RMSE', 'Parameters')] = round(error_xgboost), "n_estimator"
perform
```

```
Out[76]:
```

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10
Exponential Smoothing	781	Seasonal_periods =52
Arima	909	ARIMA(3,1,0)(2,1,0)[12]
XGBoost	714	n_estimators=1000, max_depth=5,min_child_weight=1

Model 5: FB Prophet

```
In [77]: #!pip install fbprophet
```

```
In [78]: from fbprophet import Prophet
```

```
In [79]: # Prepare data
df_train, df_test =poph2_train.reset_index(), poph2_test.reset_index()
#df_train, df_test =rabl2_train.reset_index(), rabl2_test.reset_index()
#df_train, df_test =vinb2_train.reset_index(), rabl2_test.reset_index()
df_train.columns, df_test.columns = ['ds','y'], ['ds','y']
```

```
In [80]: # Add holidays

holi = pd.DataFrame(list(holidays.UnitedKingdom(years= [2010,2011]).items()))
holi.columns=['ds','holiday']
```

```
In [81]: # Define a model
m=Prophet(holidays=holi, holidays_prior_scale=0.05) #weekly_seasonality=True, da

m.add_seasonality(name='weekly', period=7, fourier_order=3, prior_scale=5)
m.add_seasonality(name='daily', period=8, fourier_order=3, prior_scale=0.1)
m.add_seasonality(name='yearly', period=6, fourier_order=3, prior_scale=0.1)
```

Out[81]: <fbprophet.forecaster.Prophet at 0x7fa441cde5e0>

```
In [82]: m.fit(df_train)

INFO:fbprophet:Found custom seasonality named 'yearly', disabling built-in 'yearly' seasonality.
INFO:fbprophet:Found custom seasonality named 'weekly', disabling built-in 'weekly' seasonality.
INFO:fbprophet:Found custom seasonality named 'daily', disabling built-in 'daily' seasonality.
```

Out[82]: <fbprophet.forecaster.Prophet at 0x7fa441cde5e0>

```
In [83]: future=m.make_future_dataframe(periods=13, freq='D')
```


```
In [84]: prophet_pred=m.predict(future)
prophet_pred.tail()
```

Out[84]:


	ds	trend	yhat_lower	yhat_upper	trend_lower	trend_upper	Battle of the Boyne [Northern Ireland]	Battle of E [Nor Ireland]_
369	2011-12-05	318.564071	-7.675460	785.822680	318.383531	318.741389	0.0	
370	2011-12-06	320.122080	-150.624089	672.610332	319.889838	320.344679	0.0	
371	2011-12-07	321.680090	-106.786993	730.466544	321.391054	321.955181	0.0	
372	2011-12-08	323.238099	56.800815	855.213451	322.874798	323.591524	0.0	
373	2011-12-09	324.796109	91.167244	888.424066	324.362337	325.205168	0.0	

5 rows × 79 columns

```
In [85]: import matplotlib.pyplot as plt
m.plot(prophet_pred)
plt.show()
```



```
In [86]: plot_actual_predicted(poph2_test.values, prophet_pred[prophet_pred.ds>='2011-11-
```




```
In [87]: preds_prophet=prophet_pred[prophet_pred.ds>='2011-11-27']['yhat'].values
```






```
In [88]: error_prophet=np.sqrt(mean_squared_error(poph2_test.values, preds_prophet))  
print('RMSE: %.3f' % error_prophet)
```


RMSE: 790.290

```
In [89]: m.plot_components(prophet_pred)
```

Out[89]:







```
In [90]: perform.loc['Prophet', ('RMSE', 'Parameters')] = round(error_prophet), "Holidays, S
perform
```

Out[90]:	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10
Exponential Smoothing	781	Seasonal_periods =52
Arima	909	ARIMA(3,1,0)(2,1,0)[12]
XGBoost	714	n_estimators=1000, max_depth=5,min_child_weight=1
Prophet	790	Holidays, Seasonality: W, D, Y

Model 6: LSTM

```
In [91]: # Now, let's fit LSTM. Please note, no parameter tuning was performed and it is
# A model is required to learn from the series of past observations to predict t
```

```
In [92]: #!pip install tensorflow
#!pip install keras
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.preprocessing.sequence import TimeseriesGenerator
from sklearn.preprocessing import MinMaxScaler

scaler=MinMaxScaler()
```

```
In [93]: # Prepare data
train_data, test_data = pd.DataFrame(train_data), pd.DataFrame(test_data)
```

```
In [94]: # Normalize the dataset
scaler.fit(train_data)
scaled_train_data=scaler.transform(train_data)
scaled_test_data=scaler.transform(test_data)
```

```
In [95]: # Build a model
n_input = 16
n_features = 1
generator = TimeseriesGenerator(scaled_train_data, scaled_train_data, length = n
lstm_model = Sequential()
lstm_model.add(LSTM(100, activation='relu', input_shape=(n_input, n_features)))
lstm_model.add(Dense(1))
lstm_model.compile(optimizer='adam', loss='mse')
lstm_model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 100)	40800
dense (Dense)	(None, 1)	101
Total params: 40,901		
Trainable params: 40,901		
Non-trainable params: 0		

```
In [96]: # Fit a model
lstm_model.fit(generator, epochs=20, verbose=0)

losses_lstm=lstm_model.history.history['loss']
plt.figure(figsize=(12,4))
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.xticks(np.arange(0,21,1))
plt.plot(range(len(losses_lstm)), losses_lstm)
```

Out[96]: [<matplotlib.lines.Line2D at 0x7fa456c9e850>]



```
In [97]: # Predict and inverse scaling
lstm_predictions_scaled=list()
```

```

batch = scaled_train_data[-n_input:]
current_batch = batch.reshape((1, n_input, n_features))

for i in range(len(test_data)):
    lstm_pred = lstm_model.predict(current_batch)[0]
    lstm_predictions_scaled.append(lstm_pred)
    current_batch = np.append(current_batch[:, 1:, :], [[lstm_pred]], axis=1)

lstm_predictions=scaler.inverse_transform(lstm_predictions_scaled)

```

```

In [98]: # Calculate error
error_lstm=np.sqrt(mean_squared_error(test_data.values, lstm_predictions))
error_lstm


```

Out[98]: 904.9477093089544

```

In [99]: plot_actual_predicted(test_data[0], pd.DataFrame(lstm_predictions)[0], "LSTM Pre

```



```

In [100... perform.loc['LSTM', ('RMSE','Parameters')]=round(error_lstm), "LSTM: 100, act =
perform

```

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10

	RMSE	Parameters
Models Popcorn Holder		
Exponential Smoothing	781	Seasonal_periods =52
Arima	909	ARIMA(3,1,0)(2,1,0)[12]
XGBoost	714	n_estimators=1000, max_depth=5,min_child_weight=1
Prophet	790	Holidays, Seasonality: W, D, Y
LSTM	905	LSTM: 100, act = 'relu', input_shape: 16x1

In [101... *#By looking at the results so far, it is evident that seasonality doesn't have s
#My assumption is that lagged, differenciaded values and rolling mean will posit
#Let's test it out with Random Forrest by adding these features in the model's i*

Model 7: Random Forest

In [102... *# Prepare the data by differencing, adding 4 lags and rolling mean*
df_forecasting=pd.DataFrame(poph2)
df_forecasting = df_forecasting.diff()
df_forecasting.columns=['Values']
for i in range(4,0,-1):
df_forecasting['t-'+str(i)] = df_forecasting['Values'].shift(i)
df_forecasting=df_forecasting.dropna()
df_forecasting['Values_Rolling'] = df_forecasting['Values'].rolling(window = 16)
df_forecasting=df_forecasting.dropna()
df_forecasting

Out[102...

	Values	t-4	t-3	t-2	t-1	Values_Rolling
2010-12-21	-92.0	253.0	-272.0	8.0	159.0	-0.3750
2010-12-22	-75.0	-272.0	8.0	159.0	-92.0	-16.5000
2010-12-23	8.0	8.0	159.0	-92.0	-75.0	-7.5625
2010-12-24	-8.0	159.0	-92.0	-75.0	8.0	-13.2500
2010-12-25	0.0	-92.0	-75.0	8.0	-8.0	-16.2500
...
2011-12-05	296.0	-203.0	1103.0	-1369.0	113.0	25.5625
2011-12-06	-72.0	1103.0	-1369.0	113.0	296.0	7.4375
2011-12-07	31.0	-1369.0	113.0	296.0	-72.0	1.1875
2011-12-08	2738.0	113.0	296.0	-72.0	31.0	121.5000
2011-12-09	-2876.0	296.0	-72.0	31.0	2738.0	-16.6250

354 rows × 6 columns

In [103... *from sklearn.ensemble import RandomForestRegressor*
from random import seed
x=df_forecasting.iloc[:,1:]
y=df_forecasting.iloc[:,0]
x_train, x_valid = x.loc[x.index < '2011-11-27'], x.loc[x.index >= '2011-11-27']

```

y_train, y_valid = y.loc[y.index < '2011-11-27'], y.loc[y.index >= '2011-11-27']
mdl = RandomForestRegressor(n_estimators=100)
np.random.seed(55)
mdl.fit(x_train, y_train)
pred=mdl.predict(x_valid)
pred=pd.Series(pred, index=y_valid.index)

```

```

In [104... error_rf_dif=np.sqrt(mean_squared_error(y_valid, pred))
print('RMSE_dif: %.3f' % error_rf_dif)


```

RMSE_dif: 359.950

```

In [105... plot_actual_predicted(y_valid, pred, "Random Forest Fit")

```



```

In [106... # Inverse differencing and plot predicted values
converted=pd.DataFrame()
last_obs=train_data.iloc[-1][0]
converted['Conv']=np.r_[last_obs, pred[0:]].cumsum()[1:]

```

```

In [107... error_rf=np.sqrt(mean_squared_error(test_data, converted))
print('RMSE: %.3f' % error_rf)


```

RMSE: 261.448

```

In [108... plot_actual_predicted(test_data[0], converted['Conv'], "Random Forest Prediction

```



```
In [109... perform.loc['RF', ('RMSE', 'Parameters')]=round(error_rf), "Differenced: 1, Lags: 4",
perform
```

	RMSE	Parameters
Models Popcorn Holder		
Moving Average	746	Window size: 10
Exponential Smoothing	781	Seasonal_periods =52
Arima	909	ARIMA(3,1,0)(2,1,0)[12]
XGBoost	714	n_estimators=1000, max_depth=5,min_child_weight=1
Prophet	790	Holidays, Seasonality: W, D, Y
LSTM	905	LSTM: 100, act = 'relu', input_shape: 16x1
RF	261	Differenced: 1, Lags: 4 , RM Window: 16

```
In [110... #Random forrest seems to have the best fit by far. We will apply that model to t
```

```
In [111... # Therefore, the predicted order quantity for the 7 days from 11/27/2011 - 12/3/
PO_poph2=converted[0:7].sum()[0]
print('Predicted Sales Quantity Total 11/27-12/3: %.3f' % PO_poph2)
```

Predicted Sales Quantity Total 11/27-12/3: 3001.460

```
In [112... print('Actual Sales Quantity Total 11/27-12/3: %.3f' % test_data[0:7].sum()[0])
```

Actual Sales Quantity Total 11/27-12/3: 3195.000


```
In [113... # Define a function to fit RF

def RF_fit(data,lags, window):
    df_forecasting=pd.DataFrame(data)
    df_forecasting = df_forecasting.diff()
    df_forecasting.columns=['Values']
    for i in range(lags,0,-1):
        df_forecasting['t-'+str(i)] = df_forecasting['Values'].shift(i)
    df_forecasting=df_forecasting.dropna()
    df_forecasting['Values_Rolling'] = df_forecasting['Values'].rolling(window).
    df_forecasting= df_forecasting.dropna()
    x=df_forecasting.iloc[:,1:]
    y=df_forecasting.iloc[:,0]
    x_train, x_valid = x.loc[x.index < '2011-11-27'], x.loc[x.index >= '2011-11-
    y_train, y_valid = y.loc[y.index < '2011-11-27'], y.loc[y.index >= '2011-11-
    mdl = rf=RandomForestRegressor(n_estimators=100)
    np.random.seed(55)
    mdl.fit(x_train, y_train)
    pred=mdl.predict(x_valid)
    pred=pd.Series(pred, index=y_valid.index)
    error_rf_dif=np.sqrt(mean_squared_error(y_valid, pred))
    print('RMSE_dif: %.3f' % error_rf_dif)
    test_data=pd.DataFrame(data[data.index>='2011-11-27'].values)
    train_data=pd.DataFrame(data[data.index<'2011-11-27'].values)
    last_obs=train_data.iloc[-1][0]
    #Let's inverse differencing and plot predicted values
    converted=pd.DataFrame()
    converted['Conv']=np.r_[last_obs, pred[0:]].cumsum()[1:]
    error_rf=np.sqrt(mean_squared_error(test_data, converted))
    print('RMSE: %.3f' % error_rf)
    plot_actual_predicted(test_data[0], converted['Conv'], "Random Forest Predic
    #Therefore, the predicted order quantity for the 7 days from 11/27/2011 - 12
    purchase_order=converted[0:7].sum()[0]
    print('Predicted Sales Quantity Total 11/27-12/3: %.3f' % purchase_order)
    print('Actual Sales Quantity Total 11/27-12/3: %.3f' % test_data[0:7].sum()[
    return purchase_order
```

Predicting Quantities: Popcorn Holder

```
In [114... # Let's test it and ensure it works
PO_poph2=RF_fit(poph2, 4, 16)
```

RMSE_dif: 359.950
RMSE: 261.448



2000

Predicted Sales Quantity Total 11/27-12/3: 3001.460
Actual Sales Quantity Total 11/27-12/3: 3195.000

```
In [115... # While we can regress the quantities over counties and dates to come up with co
```

```
In [116... # Define a function for it
def predict_orders (data, predicted_order):
    data_train=data[(data.index<datetime.date(2011,11,27)) & (data.index>=(datet
    data_test=data[(data.index>=datetime.date(2011,11,27)) & (data.index<=dateti
    pred_countries=pd.DataFrame(data_train.groupby('Country')['Quantity'].sum()).
    actual_countries_data=pd.DataFrame(data_test.groupby('Country')['Quantity'].
    countries_data=pd.merge(pred_countries, actual_countries_data, on=['Country']
    countries_data['Ratio']=round(countries_data.Quantity_x/countries_data.Quant
    countries_data['Predictions']=round(countries_data.Ratio*predicted_order)
    countries_data.rename(columns = {'Quantity_y':'Actual'}, inplace = True)
    countries_data.rename(columns = {'Quantity_x':'Historical'}, inplace = True)
    return (countries_data) #.head(15))
```

```
In [117... predict_orders (poph, PO_poph2).head()
```


```
Out[117... Historical Actual Ratio Predictions
```

Country				
United Kingdom	12783	3083	0.98	2941.0
Italy	100	0	0.01	30.0
EIRE	92	12	0.01	30.0
France	54	0	0.00	0.0
Belgium	36	0	0.00	0.0

Predicting quantities: Vintage Doily Jumbo Bag Red

```
In [118... # Fit RF model
PO_vinb2=RF_fit(vinb2, 4, 1)
```

RMSE_dif: 521.429
RMSE: 362.303



Predicted Sales Quantity Total 11/27-12/3: 602.360
Actual Sales Quantity Total 11/27-12/3: 1851.000

```
In [119...] predict_orders(vinb, PO_vinb2).head()
```


```
Out[119...]      Historical  Actual  Ratio  Predictions
Country
United Kingdom    2125    1830   0.94         566.0
Portugal           40         0   0.02          12.0
France            35        10   0.02          12.0
Germany            20        10   0.01           6.0
Finland            20         0   0.01           6.0
```

Predicting Quantities: Rabbit Night Light

```
In [120...] # Fit the Random Forest to Rabbit Night Light data
```

```
In [121...] PO_rabl2=RF_fit(rabl2, 4, 1)
```

RMSE_dif: 485.418
RMSE: 337.100



Predicted Sales Quantity Total 11/27-12/3: 3906.450
Actual Sales Quantity Total 11/27-12/3: 4588.000

```
In [122...] predict_orders(rabl, PO_rabl2).head()
```

```
Out[122...]      Historical  Actual  Ratio  Predictions
Country
United Kingdom    8784    1985    0.51    1992.0
Netherlands       2616         0    0.15     586.0
France            2326     383    0.14     547.0
Australia          1632         0    0.10     391.0
Japan             1080    2040    0.06     234.0
```

```
In [123...] # Japan's order is obviously underpredicted
```

```
In [124...] rabl[rabl.Country=="Japan"] #It was unusual order for that country
```

```
Out[124...]      InvoiceNo  StockCode  Description  Quantity  UnitPrice  CustomerID  Country
InvoiceDate
```

InvoiceDate	InvoiceNo	StockCode	Description	Quantity	UnitPrice	CustomerID	Country
2011-06-22	557670	23084	RABBIT NIGHT LIGHT	288	1.79	12798.0	Japan
2011-10-26	572869	23084	RABBIT NIGHT LIGHT	960	1.79	12798.0	Japan
2011-11-17	576923	23084	RABBIT NIGHT LIGHT	120	1.79	12753.0	Japan
2011-11-29	579498	23084	RABBIT NIGHT LIGHT	2040	1.79	12798.0	Japan
2011-12-06	C580832	23084	RABBIT NIGHT LIGHT	-7	1.79	12753.0	Japan

```
In [125... # Add predicted orders to top3 dataframe
orders={'Rabbit_Night_Light': round(PO_rabl2), 'Popcorn_Holder': round(PO_poph2)}
```

```
In [126... top3['Description']=[x for x in pd.DataFrame(orders.items())[0]]
```

```
In [127... top3['Predicted_Order_ML']=[x for x in pd.DataFrame(orders.items())[1]]
```

```
In [128... top3
```

StockCode	Quantity	Description	Predicted_Order_ML
23084	4588	Rabbit_Night_Light	3906
22197	3195	Popcorn_Holder	3001
23582	1851	Vintage_Doily_Jumbo_Bag	602

```
In [129... # A workable solution to address understock is to add a safety stock:
# Safety stock ss = std(delivery_lead_time) * importance_factor
# Since we don't have that data, we can substitute:
# ss = std(daily_quantity) * (days_of_prediction) * (model_confidence_factor)
# Therefore, our prediction could be the following:
# Predicted_Order = Predicted_Order_ML + ss
```

```
In [130... model_confidence_factor=[0.3, 0.1, 0.9]
stdevs=[np.std(rabl2.values)*6,np.std(poph2)*6,np.std(vinb2)*6]
```

```
In [131... ss = [round(a * b) for a, b in zip(model_confidence_factor, stdevs)]
print(ss)
```

```
[645, 224, 1201]
```

```
In [132... top3.insert(2, "Safety_Stock", ss)
```

```
In [133... top3['Final_Order']=top3.Safety_Stock+top3.Predicted_Order_ML
```

```
In [134... top3
```

Out[134...

	Quantity	Description	Safety_Stock	Predicted_Order_ML	Final_Order
StockCode					
23084	4588	Rabbit_Night_Light	645	3906	4551
22197	3195	Popcorn_Holder	224	3001	3225
23582	1851	Vintage_Doily_Jumbo_Bag	1201	602	1803

In [135...

```
# Aggregate the results
rabbit_light=predict_orders(rabl, top3.Final_Order.values[0])
rabbit_light.head()
```

Out[135...

	Historical	Actual	Ratio	Predictions
Country				
United Kingdom	8784	1985	0.51	2321.0
Netherlands	2616	0	0.15	683.0
France	2326	383	0.14	637.0
Australia	1632	0	0.10	455.0
Japan	1080	2040	0.06	273.0

In [136...

```
popcorn_holder=predict_orders(poph, top3.Final_Order.values[1])
popcorn_holder.head()
```

Out[136...

	Historical	Actual	Ratio	Predictions
Country				
United Kingdom	12783	3083	0.98	3160.0
Italy	100	0	0.01	32.0
EIRE	92	12	0.01	32.0
France	54	0	0.00	0.0
Belgium	36	0	0.00	0.0

In [137...

```
vintage_bag=predict_orders(vinb, top3.Final_Order.values[2])
vintage_bag.head()
```

Out[137...

	Historical	Actual	Ratio	Predictions
Country				
United Kingdom	2125	1830	0.94	1695.0
Portugal	40	0	0.02	36.0
France	35	10	0.02	36.0
Germany	20	10	0.01	18.0
Finland	20	0	0.01	18.0

```
step1=pd.merge(rabbit_light, popcorn_holder, on = 'Country', how='outer')
```

```
In [138...
```

```
In [139... final_data=pd.merge(step1, vintage_bag, on = 'Country', how='outer')
```

```
In [140... final_data.columns=['Hist_rabbbbit_light', 'Act_rabbit_light', 'R_rabbit_light',  
                        'Preds_popcorn_holder', 'Hist_vintage_bag', 'Act_vintage_bag'
```

```
In [141... final_data
```

Out[141...

	Hist_rabbbbit_light	Act_rabbit_light	R_rabbit_light	Preds_rabbit_light	Hist_popcorn_hc
Country					
United Kingdom	8784	1985	0.51	2321.0	1:
Netherlands	2616	0	0.15	683.0	
France	2326	383	0.14	637.0	
Australia	1632	0	0.10	455.0	
Japan	1080	2040	0.06	273.0	
Germany	192	72	0.01	46.0	
Belgium	108	0	0.01	46.0	
Finland	96	48	0.01	46.0	
Sweden	84	0	0.00	0.0	
EIRE	48	0	0.00	0.0	
Iceland	48	0	0.00	0.0	
Italy	48	0	0.00	0.0	
Denmark	24	12	0.00	0.0	
Portugal	18	48	0.00	0.0	
Unspecified	12	0	0.00	0.0	
Norway	12	0	0.00	0.0	
Switzerland	12	0	0.00	0.0	
Spain	6	0	0.00	0.0	
Lithuania	0	0	0.00	0.0	
USA	0	0	0.00	0.0	
Austria	0	0	0.00	0.0	
Bahrain	0	0	0.00	0.0	
United Arab Emirates	0	0	0.00	0.0	
Brazil	0	0	0.00	0.0	
Canada	0	0	0.00	0.0	
Channel Islands	0	0	0.00	0.0	

	Hist_rabbbbit_light	Act_rabbit_light	R_rabbit_light	Preds_rabbit_light	Hist_popcorn_hc
Country					
Cyprus	0	0	0.00	0.0	
Czech Republic	0	0	0.00	0.0	
European Community	0	0	0.00	0.0	
Lebanon	0	0	0.00	0.0	
Singapore	0	0	0.00	0.0	
Saudi Arabia	0	0	0.00	0.0	
Greece	0	0	0.00	0.0	
Hong Kong	0	0	0.00	0.0	
RSA	0	0	0.00	0.0	
Poland	0	0	0.00	0.0	
Malta	0	0	0.00	0.0	
Israel	0	0	0.00	0.0	

```
In [142... final_data.to_csv(r'Tatsiana_Sokalava}_result.csv', index = True, header=True)
```

Thank you!